

# Evaluation Framework Design

AI-Driven Compiler Optimization System

*Week 2 Deliverable*

February 2026

## Executive Summary

This document defines a comprehensive evaluation framework for measuring the success of the AI-Driven Compiler Optimization System. The framework establishes quantitative and qualitative metrics, defines benchmark datasets, specifies testing methodologies, and outlines comparison protocols against baseline systems.

The evaluation framework is designed to ensure rigorous, reproducible, and comprehensive assessment of the system across five key dimensions:

- Correctness: Semantic equivalence and bug-free transformations
- Performance: Execution speed improvements and optimization quality
- Security: Absence of newly introduced vulnerabilities
- Explainability: Quality and completeness of reasoning
- Usability: Developer experience and integration ease

## **Table of Contents**

1. Evaluation Objectives and Success Criteria
2. Evaluation Metrics
  - 2.1 Correctness Metrics
  - 2.2 Performance Metrics
  - 2.3 Security Metrics
  - 2.4 Explainability Metrics
  - 2.5 Usability Metrics
3. Benchmark Datasets and Test Suites
4. Baseline Comparison Systems
5. Testing Methodology
6. Statistical Analysis Methods
7. Evaluation Timeline and Milestones
8. Reporting and Documentation

# 1. Evaluation Objectives and Success Criteria

The primary objective of this evaluation framework is to rigorously assess whether the AI-Driven Compiler Optimization System achieves its stated goals while maintaining formal correctness and security guarantees.

## 1.1 Primary Success Criteria

Criterion	Target	Priority
Correctness Rate	≥95%	Critical
Performance Improvement	≥10% speedup in ≥30% of cases	High
Optimization Detection	30% more opportunities than static analyzers	High
Security Guarantee	Zero new vulnerabilities	Critical
Explainability Coverage	100% of suggestions	High

## 1.2 Secondary Success Criteria

- Developer Trust: Average rating ≥4.0/5.0 on explanation clarity (user study with 10+ developers)
- Integration Ease: Successful integration with 3+ real-world projects without code modifications
- Performance Overhead: Total optimization time <10x original compilation time
- Failure Handling: Graceful degradation with clear error messages in 100% of failure cases

## 2. Evaluation Metrics

### 2.1 Correctness Metrics

Correctness is the most critical evaluation dimension, as incorrect optimizations render the system unusable regardless of performance improvements.

#### 2.1.1 Semantic Equivalence

Metric	Measurement Method	Target
<b>Formal Verification Pass Rate</b>	Z3 SMT solver equivalence proofs	100%
<b>Differential Testing Pass Rate</b>	Automated test generation with 100+ inputs per function	≥90%
<b>Manual Code Review Validation</b>	Expert review of 100 randomly sampled optimizations	≥90% approval

#### 2.1.2 Bug Introduction Rate

- False Positive Rate: Percentage of suggested optimizations that introduce bugs (Target: <10%)
- Regression Detection: Number of previously passing tests that fail after optimization (Target: 0)
- Edge Case Handling: Correct behavior on corner cases (null pointers, integer overflow, etc.) (Target: 100%)

## 2.2 Performance Metrics

### 2.2.1 Execution Speed Improvement

Metric	Measurement Method	Target
<b>Average Speedup</b>	Google Benchmark on optimized vs baseline(optional)	Report mean, median, std dev
<b>Percentage of Improved Cases</b>	Functions with ≥20% speedup	≥10%
<b>Best Case Improvement</b>	Maximum speedup observed	Report top 5%
<b>Worst Case Analysis</b>	Functions with slowdown	<5% slower, document reasons
<b>Asymptotic Improvement</b>	Algorithmic complexity improvements ( $O(n^2) \rightarrow O(n \log n)$ )	Report all cases

#### 2.2.2 Resource Utilization

- Memory Usage: Peak memory consumption (optimized vs baseline)
- Code Size: Binary size comparison (optimized vs baseline)

- Energy Consumption: Estimated energy usage based on execution time and CPU frequency

### 2.2.3 Optimization Coverage

- Optimization Opportunities Identified: Total number of potential optimizations detected
- Optimization Opportunities Applied: Number of successfully applied optimizations
- Coverage by Category: Algorithm selection, data structure changes, loop optimizations, etc.

## 2.3 Security Metrics

Metric	Measurement Method	Target
New Vulnerabilities Introduced	Static analysis (Clang, cppcheck) + manual audit	0
Buffer Overflow Detection	AddressSanitizer on all optimized code	0 new issues
Race Condition Detection	ThreadSanitizer on concurrent code	0 new issues
Use-After-Free Detection	AddressSanitizer + manual review	0 new issues
Side-Channel Vulnerability	Manual audit of crypto-related optimizations	0 new issues
Security Test Suite Pass Rate	Juliet Test Suite (CWE patterns)	100%

## 2.4 Explainability Metrics

### 2.4.1 Reasoning Quality

Metric	Measurement Method	Target
Reasoning Coverage	Percentage of optimizations with complete reasoning chains	100%
Citation Completeness	Percentage of claims with code line references	$\geq 90\%$
Reasoning Step Validity	Logical consistency checking by verification agent	$\geq 95\%$
Performance Prediction Accuracy	Predicted vs actual speedup correlation	$r^2 \geq 0.7$
Trade-off Explanation	Percentage of suggestions explaining trade-offs	100%

### 2.4.2 User Comprehension

User Study Questionnaire (10+ developers, Likert scale 1-5):

- Clarity: "The explanation clearly describes what optimization is being applied" (Target:  $\geq 4.0$ )
- Rationale: "I understand why this optimization improves performance" (Target:  $\geq 4.0$ )
- Trust: "I trust this optimization is correct based on the explanation" (Target:  $\geq 4.0$ )
- Actionability: "I can make an informed decision to accept/reject" (Target:  $\geq 4.2$ )

## 2.5 Usability Metrics

### 2.5.1 Integration Metrics

- Setup Time: Time to integrate into existing project (Target: <30 minutes)
- Build System Compatibility: Works with Make, CMake, Ninja (Target: 100%)
- Breaking Changes: Number of code modifications required (Target: 0)

### 2.5.2 Performance Overhead

- Analysis Time: Time to analyze codebase (Target: <2x compilation time)
- Verification Time: Time for formal verification (Target: <5x analysis time)
- Total Overhead: End-to-end time (Target: <10x baseline compilation)

## 3. Benchmark Datasets and Test Suites

### 3.1 Correctness Benchmarks

Dataset	Size	Purpose	Source
Algorithm Implementations	20 functions	Sorting, searching, graph algorithms	GitHub, LeetCode
Known Optimizations	15 functions	Ground truth optimized versions	Expert-created
Edge Cases	10 functions	Null pointers, overflows, boundary conditions	Custom-generated
Real-World Code	10 functions	Production code from open-source projects	Apache/MIT licensed repos
Adversarial Examples	15 functions	Deliberately complex/tricky code	Custom-created

### 3.2 Performance Benchmarks(optional)

Benchmark Suite	Programs	Domain	Metrics
SPEC CPU 2017	20 programs	General computation	Execution time, memory
Scientific Computing	15 programs	Matrix operations, numerical methods	FLOPs, accuracy
Data Processing	10 programs	Parsing, transformation	Throughput, latency
Web Services	10 programs	HTTP handling, JSON processing	Requests/sec
Custom Microbenchmarks	50 functions	Specific optimization patterns	Targeted metrics

### 3.3 Security Test Suites

Test Suite	Test Cases	Vulnerability Types	Coverage
Custom Vulnerability Tests	100	Optimization-specific security issues	Targeted

## 4. Baseline Comparison Systems

The AI-Driven Compiler Optimization System will be compared against the following baselines to demonstrate its unique value and capabilities:

### 4.1 Traditional Compilers(post integration-optional)

Compiler	Version	Optimization Levels	Comparison Focus
LLVM/Clang	16.x	-O1, -O2, -O3,	IR-level optimizations
GCC	12.x	-O1, -O2, -O3,	Traditional optimizations

### 4.2 AI Code Assistants

Tool	Capability	Comparison Aspect
ChatGPT (GPT-4)	Code analysis & refactoring	Semantic understanding

### 4.3 Static Analysis Tools

Tool	Analysis Type	Comparison Metric
cppcheck	C/C++ static analysis	Performance warnings

### 4.4 Comparison Methodology

**Same Codebase Testing:** All systems tested on identical benchmark datasets

**Controlled Environment:** Same hardware, OS version, system load

**Multiple Runs:** Minimum 10 runs per benchmark, report mean and confidence intervals

**Fair Comparison:** Best settings for each baseline (e.g., -O3 for GCC/LLVM)

**Documented Differences:** Clear attribution of improvements to specific techniques

## 5 Statistical Analysis Methods

### 5.1 Performance Analysis

- Descriptive Statistics: Mean, median, standard deviation, min, max for all metrics
- Confidence Intervals: 95% CI for mean speedup and other continuous metrics

### 5.2 Categorical Analysis

- Optimization Type Distribution: Frequency analysis by optimization category
- Failure Mode Analysis: Categorization and frequency of failure types

## 6. Evaluation Timeline and Milestones

Week	Evaluation Activity	Deliverable
4	Dataset Collection	Benchmark datasets prepared (50+ test cases)
8	Preliminary Testing	Core functionality tested
10	Integration Testing	End-to-end pipeline tested
12	Failure Analysis	Failure taxonomy and mitigation strategies
13	Comprehensive Evaluation	Full benchmark suite execution
13	User Study	Developer feedback collection
14	Final Analysis & Reporting	Complete evaluation report with statistical analysis

## **8. Reporting and Documentation**

### **8.1 Evaluation Report Structure**

- Executive Summary: Key findings and overall assessment
- Methodology: Detailed description of evaluation procedures
- Results by Metric: Quantitative results for each metric with statistical analysis
- Baseline Comparisons: Side-by-side comparisons with tables and graphs
- Case Studies: Detailed analysis of 10-20 representative examples
- Failure Analysis: Categorized failures with root causes
- Discussion: Interpretation of results, limitations, future work
- Appendices: Complete data tables, benchmark code, statistical tests

### **8.2 Visualization Requirements**

- Performance Graphs: Box plots, violin plots for speedup distribution
- Correctness Dashboard: Pass/fail rates across test categories
- Comparison Charts: Bar charts comparing against baselines
- Heat Maps: Optimization types vs success rates
- Scatter Plots: Predicted vs actual performance

### **8.3 Data Archival**

- Raw Data: All benchmark results stored in CSV/JSON format
- Test Artifacts: Original and optimized code for all test cases
- Analysis Scripts: Reproducible statistical analysis code
- Version Control: Git repository with tagged releases for each evaluation

## Conclusion

This evaluation framework provides a rigorous, comprehensive, and reproducible methodology for assessing the AI-Driven Compiler Optimization System across all critical dimensions: correctness, performance, security, explainability, and usability.

The framework is designed to:

- Ensure Reliability: Multi-layered verification ensures formal correctness guarantees
- Enable Fair Comparison: Standardized benchmarks and baselines allow objective assessment
- Support Continuous Improvement: Detailed failure analysis guides system enhancements
- Maintain Transparency: Comprehensive documentation ensures reproducibility

By adhering to this framework, the project will generate credible evidence of the system's capabilities and limitations, supporting both academic publication and practical deployment decisions.

## Appendix A: Evaluation Checklist

Evaluation Component	Status (✓/✗)
Benchmark datasets collected and documented	
Baseline systems configured and tested	
Automated testing pipeline implemented	
Manual review protocol established	
User study designed and participants recruited	
Statistical analysis scripts prepared	
Correctness metrics ≥95% achieved	
Performance targets met (≥20% in ≥30% cases)	
Zero new security vulnerabilities confirmed	
Evaluation report drafted and reviewed	